

Kubernetes Introduction

Sandor Guba - Infrastructure Engineer at IBM Budapest Labs

\$: whoami

- Infrastructure Engineer at IBM Budapest Labs
- Working with Linux servers more than 10 years
- Main topics are Virtualization and Containerization both in private and public cloud
- And a little bit of networking ;)

What is this presentation about?

- The beginning of containerization decades ago...
- Some basic concepts
- And Kubernetes of course
 - Architecture
 - Quick walkthrough
 - Some basic example
 - Useful links and tools



First there were Bare Metal

- At the beginning everything runs on physical server
 - It worked
 - Gets troublesome with the growing number of servers
 - Lots of different tools and unique solutions hard to maintain
- Later with the evolution of virtualization “everything” runs inside a VM
 - Optimized workloads
 - A bit more secure
 - Divide workload into smaller units
 - Operating System is a considerable overhead
- Next step?

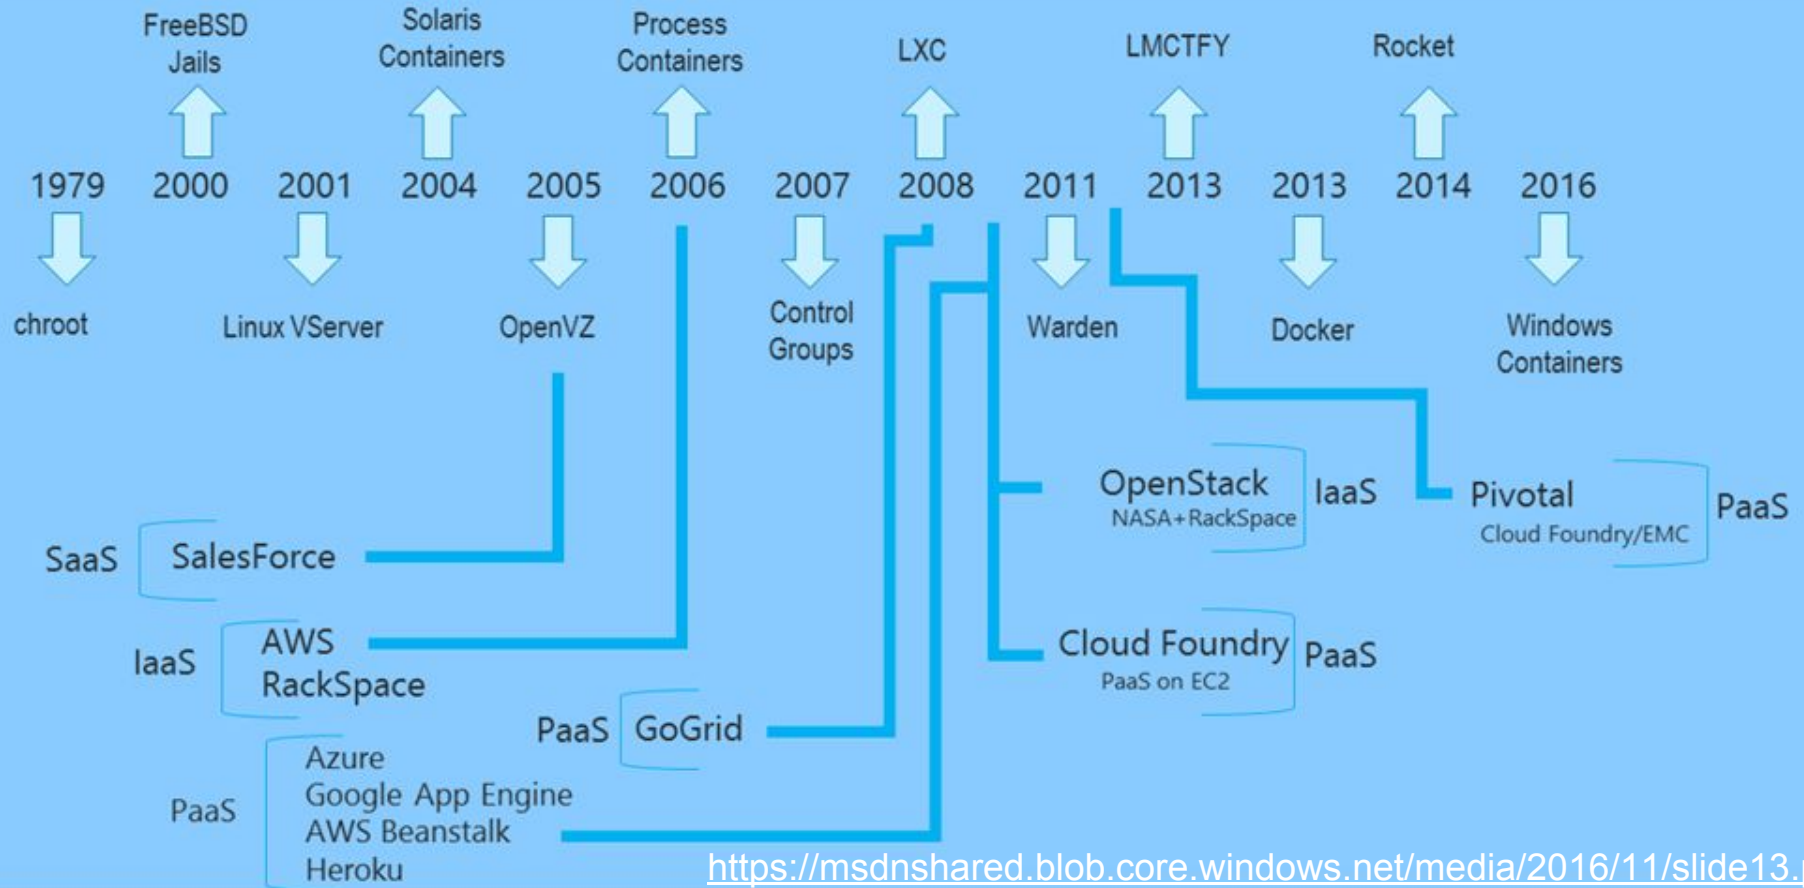


Next Step is Containerization

- The goal is a virtual machine like separation
- Without the OS overhead
- Immutable deployment
 - Production like behaviour during development
- Lots of problem to solve (mainly in Linux kernel)

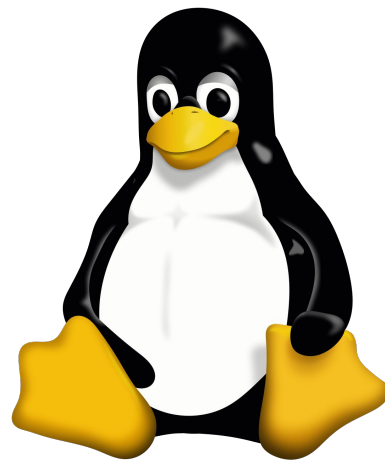


Surprising to most is that containers have been around for decades

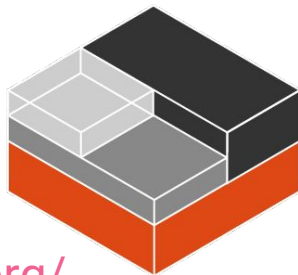


The beginning of the road to Cloud Native

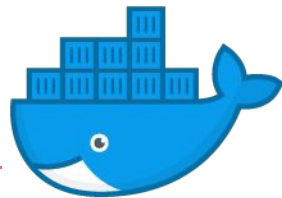
- Namespaces
 - Mnt, pid, net, ipc, hostname, user ids (isolate process environments)
- Cgroups
 - Cpu, memory, disk, i/o (limit resource usages)
- AppArmor, SELinux
 - security/access control
- Seccomp
 - Compute isolation
- Chroot
 - Filesystem isolation



Container runtimes



- LXC, LXD, CGManager, LXCFS - <https://linuxcontainers.org/>
- Docker - <https://www.docker.com/>
- Rkt - <https://coreos.com/rkt/docs/latest/>
- Open Container Initiative - <https://www.opencontainers.org/>
- Cri-o (OCI implementation)
<https://github.com/kubernetes-incubator/cri-o>



docker

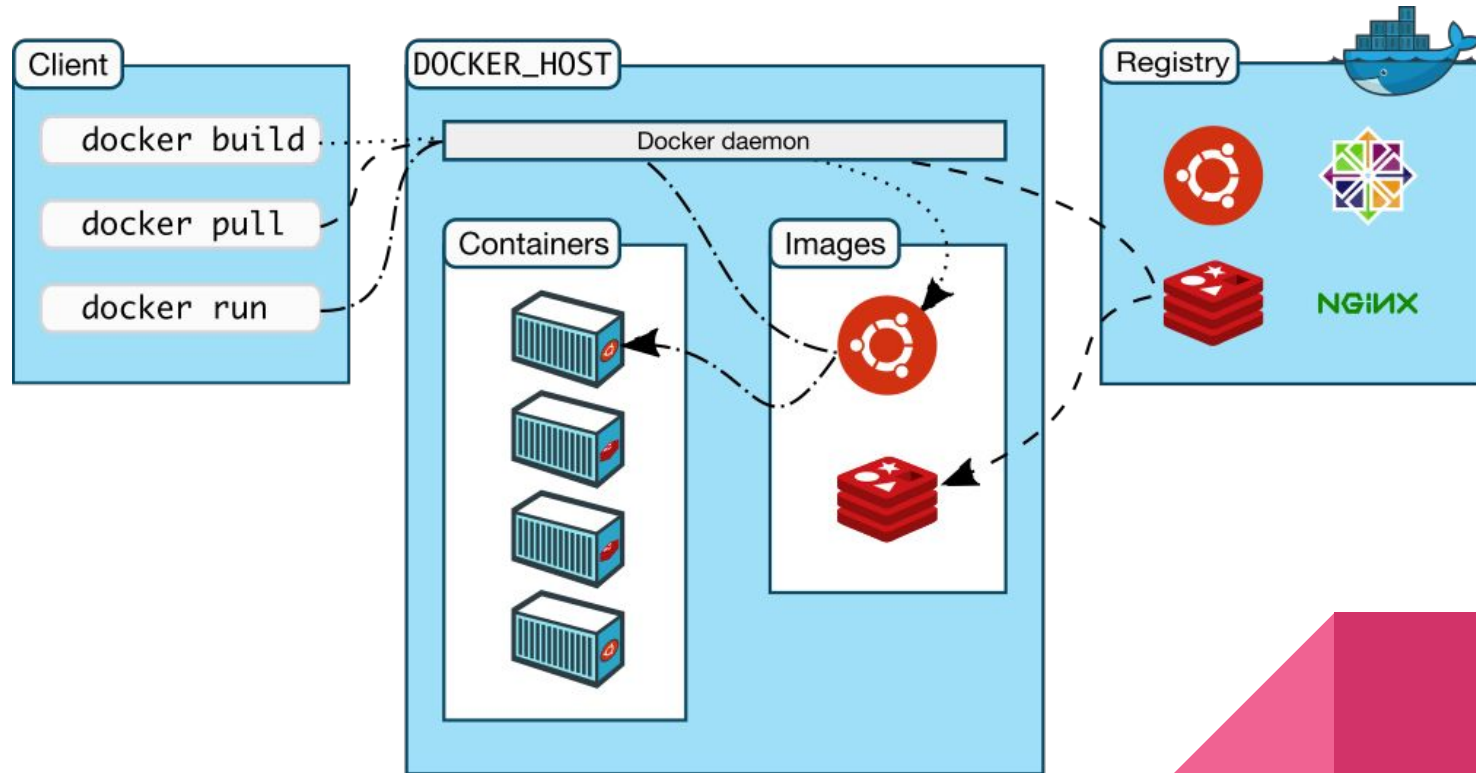


cri-o



rkt

Docker ecosystem



Here we go to Kubernetes



What is Kubernetes?

An Open Source software to automating

- Deployment
- Scaling
- Management

Of Containerized applications.

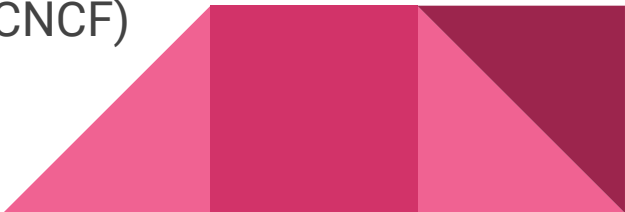


Long story short

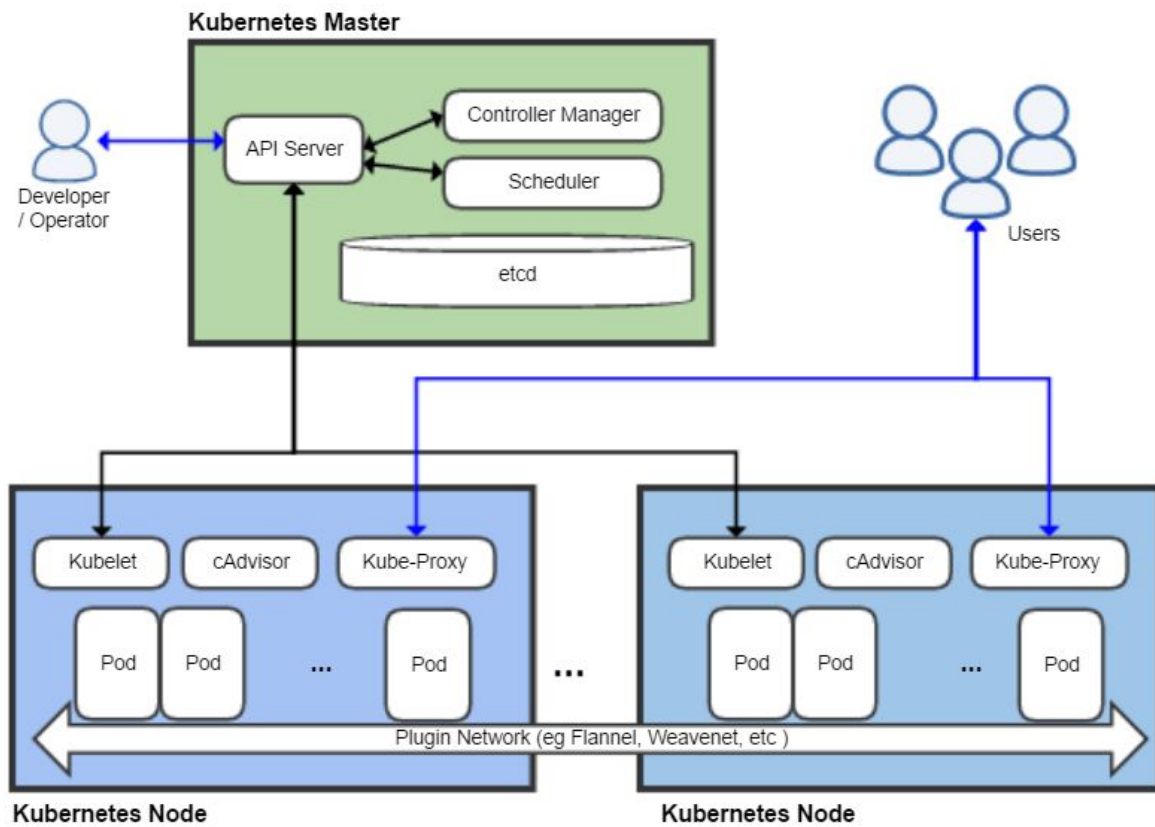
- Kubernetes is based on Google Borg
- OpenSource
- Rapidly evolving
 - It was first announced by Google in mid-2014
 - Kubernetes v1.0 was released on July 21, 2015
 - As today (Nov 21, 2017) v1.9 is in progress
- Google partnered with the [Linux Foundation](#) to form the [Cloud Native Computing Foundation](#)(CNCF)^[12] and offered Kubernetes as a seed technology.
- Big Community



Kubernetes community

- Really active open source project
 - 29k stars, 1400+ contributors, ~60k commits
 - Apache 2 licensed
 - Written in Go
 - Hosted by the Cloud Native Computing Foundation (CNCF)
- 

Architecture



POD

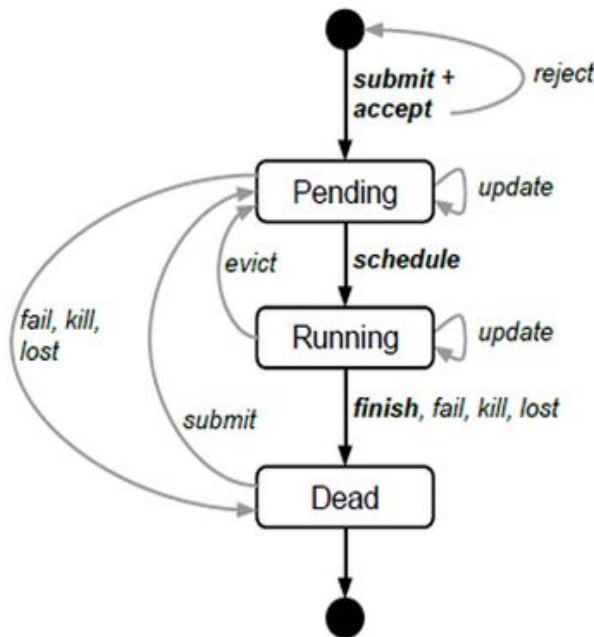
- “Smallest deployable units of a computing”
- It consist 1 or more container
- Containers shares IP address, ports and namespace
 - Localhost
 - Shared memory
- Ephemeral



Lifecycle of a POD

- Create -> Running -> Delete
- Readiness and Liveness probe
 - Exec
 - TCP
 - HTTP
- RestartPolicy
 - Always
 - OnFailure
 - Never

```
livenessProbe:  
  httpGet:  
    # scheme: HTTPS  
    path: /healthz  
    port: 8080  
    httpHeaders:  
      - name: X-Custom-Header  
        value: Awesome  
    initialDelaySeconds: 15  
    timeoutSeconds: 1  
  name: liveness
```



Deployment and ReplicaSet

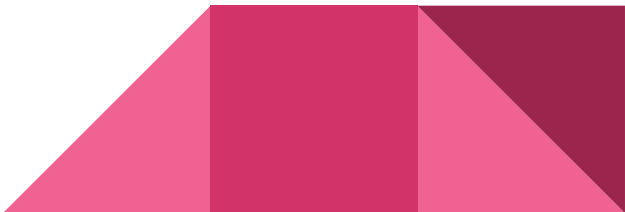
- Declarative representation of Pods and Replica Sets
- Rollout and Rollback different versions
- Manage scaling of Pods

```
apiVersion: apps/v1beta2
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

Service

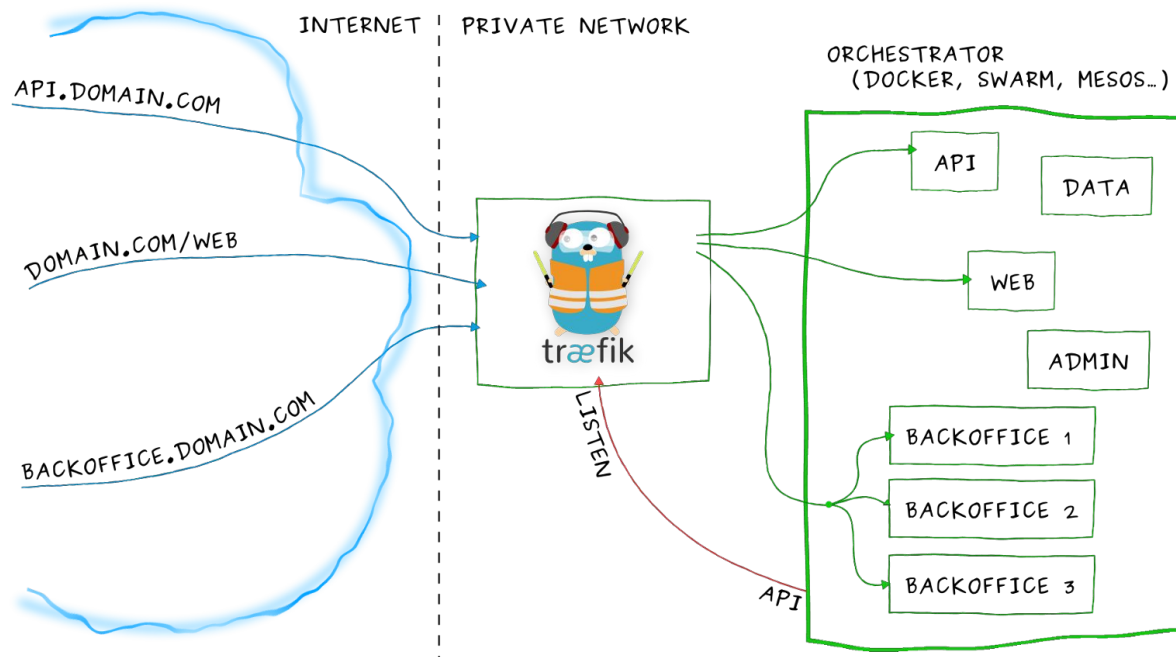
- Abstraction layer that defines a set of Pods
- Service discovery (DNS based)
- Cloud support (AWS, Google, ...)
- Multiple mode:
 - NodePort (HostPort)
 - Headless (Direkt POD addresses)
 - LoadBalancer (Internal or External with VIP)

```
kind: Service
apiVersion: v1
metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```



Ingress

- L7 Load Balancer
- Routing based on
 - Host
 - Path
 - Custom rules
- Implementations
 - nginx
 - traefik



NGINX

Namespaces

- Virtual clusters called namespaces
- Resources can be separated by namespaces
- Quotas are applied to namespaces
 - Requests cpu/memory (for scheduling)
 - Limits cpu/memory (for hard usage limit SIGKILL)
- Service names separated by namespaces by design
 - `<service-name>.<namespace-name>.svc.cluster.local`



Jobs and Cron Jobs

- As the name tells these are Pods for a specific task
- Or periodic task
- Example
 - ElasticSearch Curator
 - Backup scripts

```
apiVersion: batch/v1
kind: Job
metadata:
  name: pi
spec:
  template:
    metadata:
      name: pi
    spec:
      containers:
        - name: pi
          image: perl
          command: ["perl", "-Mbignum=bpi",
"-wle", "print bpi(2000)"]
          restartPolicy: Never
      backoffLimit: 4
```

StatefulSet

- Guaranties the order and uniqueness of a set of Pods
- Fix naming `$(statefulset name)-$(ordinal)`
- Example
 - MySQL cluster
 - Kafka
 - Cassandra

DaemonSet

- Ensure that Pod is running on all or a set of nodes
- Example
 - Fluentd
 - Kube-metrics



+1 Kubectl and Dashboard

- Your best friend communicating with kubernetes
- 1:1 mapping to the REST API
- Simple, deterministic interface
- Useful commands
- Basic proxy, port-forward functions

Examples

```
$: kubectl get nodes
```

```
$: kubectl get pod
```

```
$: kubectl get services
```

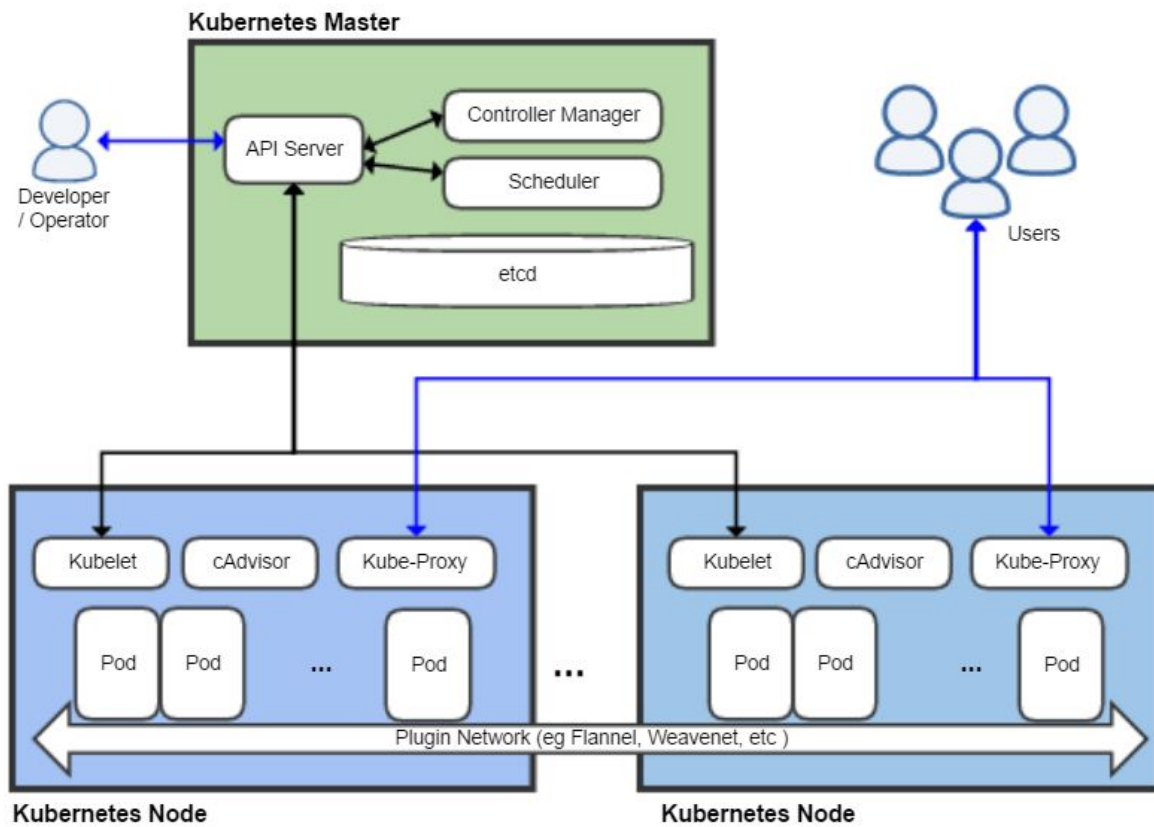
```
$: kubectl create deployment -f nginx.yaml
```

```
$: kubectl port-forward <pod_name> 8080:8080
```

```
$: kubectl proxy
```



Architecture



High Level Architecture

Kubernetes **Master** Components

- API Server
- Etcd
- Scheduler
- Controller Manager

Kubernetes **Node** Components

- Kubelet
- Kube-proxy

Kubernetes **Concepts**:

- POD
- Deployment
- Service
- Ingress

+1 Kubectl



API server

- REST operations for the resources managed by Kubernetes
- Stateless API backed by etcd
- `curl --cert myuser.pem --key myuser-key.pem --cacert /path/to/ca.pem https://apiserver:6443/api/v1/pods`
- API reference <https://kubernetes.io/docs/api-reference/v1.8/>
- Kubectl command line tool



etcd



- Distributed and reliable key-value store
- Secure with TLS
- Fast (<https://coreos.com/blog/performance-of-etcd.html>)
- written in Go and uses the Raft consensus algorithm



Scheduler

- Topology aware, workload specific functions
- Schedule based on
 - Requirements
 - Limits
 - Affinity
 - Policy
- Pluggable (custom schedulers)



Controller Manager

- Embeds the core control loops
- Watches the state of the cluster and makes changes
- Handling
 - replication controller
 - endpoints controller
 - namespace controller
 - serviceaccounts controller.
 - +1 cloud provider controller




Kubelet and Kube Proxy

Kubelet

- It is the primary node agent
- Handles POD spec and runs pods

Kube Proxy

- Network proxy that handles TCP, UDP forwarding across a set of backends
 - It handles service IP as well (usually backed by DNS)
- 

Networking

- CNI - Container Network Interface (pluggable)
- Calico (native L3)
- Flannel (tunnel)
- Weave (tunnel)
- And there are many many more...



And many many more...

- Role Based Access Control (RBAC)
 - Control what users have access to what objects
- Multiple Schedulers
- Flexible Scheduling Constraints
 - Affinity, anti-affinity, taints, tolerations
- Automatic Cluster Scaling
- K8s publishes signals that allow external services to scale the cluster automatically.
- Cloud Provider Integration
 - GCP, AWS, Azure, OpenStack, vSphere
- Network Policy
- Network ingress policy



Okay where should I start?

Minikube - <https://github.com/kubernetes/minikube>

Katacoda - <https://www.katacoda.com/>

Deploying K8s - <https://github.com/kelseyhightower/kubernetes-the-hard-way>

Introduction to K8s - <https://www.edx.org/course/introduction-kubernetes-linuxfoundationx-lfs158x>

Kubernetes - <https://kubernetes.io/>

Helm - <https://helm.sh/>





Questions?